

CO661 - Theory and Practice of Concurrency

Notes on the Calculus of Communicating Systems (CCS)

Last updated on April 9, 2019

In this module, we examine both the theory and practice of concurrency. The theoretical side brings rigour to our reasoning and understanding of key concepts in concurrent programming. In this module, we utilise CCS, the Calculus of Communicating Systems,¹ as a model for understanding concurrency. CCS is essentially a small programming language (a “*calculus*”) which is carefully and formally defined. This allows us to be precise about the meaning of concurrent programs. CCS is often described as a *process calculus* or *process algebra* since it is a calculus where “processes” are the central unit of organisation: its programs are processes, which are made up of further processes.

These notes are open source.² If you spot any mistakes or errors or things that need clarifying, please fork and send a pull request.

Contents

1 Syntax	1
2 Semantics	2
2.1 Collected rules	6
3 Modelling locks, semaphores, and race conditions	7
4 Structural congruence	9
5 Bisimulation	10

1 Syntax

CCS processes comprise *actions*, which we use to abstract the notion of a *command* in a program or a particular function call or message. For example, the following CCS program describes a process that performs an (abstract) action a followed by an (abstract) action b and then does nothing:

$$a.b.\mathbf{0}$$

Actions sequenced via the operator “.” which is called the *prefixing operator* since it is used to prefix a process with an action. The process $\mathbf{0}$ is the inactive process.

The full syntax of CCS is defined inductively by the following BNF grammar:

$P, Q ::=$	$\alpha.P$	<i>prefixing</i>
	$P + Q$	<i>non-deterministic choice</i>
	$(P \mid Q)$	<i>parallel composition</i>
	$\nu a.P$	<i>name restriction</i>
	$\mathbf{0}$	<i>inactive process</i>
	$P[b/a]$	<i>relabelling</i>
	A	<i>process identifiers</i>

where $\alpha ::= a \mid b \mid \dots \bar{a} \mid \bar{b} \mid \dots$ ranges over *names* (of actions) where names with an overline are called *co-names* (or *dual names*). We use P and Q to range over processes, i.e., when we want to refer

¹Robin Milner: *A Calculus of Communicating Systems*, Springer Verlag, ISBN 0-387-10235-3. 1980.

²<https://github.com/dorchard/co661-notes.git>

to some process in an abstract sense. We use A to range over identifiers for processes (i.e., concrete names given to processes, see below on “Process definitions”).

The following gives a few examples to give some intuition, but in the next section we will look at the concrete semantics of processes which will make precise what each part of the syntax means.

- $a.b.P$ – does action a then action b then continues to behave as the process P ;
- $a.0 + b.0$ – chooses non-deterministically between doing action a or doing action b .
- $(a.b.0 \mid c.d.0)$ – in parallel, one process does action a then b then stops (inactive process), whilst the other does c then d then stops. What is observed is an *interleaving* of these actions, e.g., one possibility is that whole process does a then c then b then d . Another possibility is c then d then a then b . The point is that we may get any interleaving of the actions of the subprocesses.
- $\nu a.(a.P \mid \bar{a}.Q)$ – this is a process where no a action can be observed externally, because of the name restriction νa . The inner process $(a.P \mid \bar{a}.Q)$ can do a *handshake* between the action a and its dual \bar{a} which can happen together leaving a process $P \mid Q$.

Concurrency Workbench Note that the Concurrency Workbench tool (CAAL), which we use in this course, has a slightly different concrete syntax, with the above constructs represented as follows (in order):

$$P ::= a.P \mid 'a.P \mid P+Q \mid (P \mid Q) \mid P \setminus \{a1, \dots, an\} \mid 0 \mid P[b/a] \mid A$$

Process definitions We will defined processes by recursive definitions. For example, the following defines a process named A :

$$A \stackrel{def}{=} a.b.A$$

The process is recursively defined, and performs actions a then b continuously.

Processes can be mutually recursive. For example, we could define a process that performs a then b actions continuously via two definitions:

$$A \stackrel{def}{=} a.B \qquad B \stackrel{def}{=} b.A$$

Here A performs a then b actions continuously whilst B performs b then a actions continuously (they have different starting points).

2 Semantics

We describe the semantics (the meaning) of CCS processes via a *labelled-transitions system* which effectively describes small “steps” of evaluation for CCS processes along with actions which are made observable by this evaluation. We define these small steps (called *reductions* or *transitions*) by a relation between processes,³ annotated with a label, of the form:

$$P \xrightarrow{l} Q$$

which describes that a process P can reduce (*evaluate*) to a process Q by doing some action described by the label l . This is a *labelled transition*. Labels are as defined as:

$$l ::= \alpha \mid \tau$$

i.e., they are either names of actions or they are a special label τ called a *silent action*.

Since the syntax of processes is itself inductively defined (i.e. processes can contain processes) so the definition of \rightarrow is inductively defined. We leverage the *inference rule* style, with inductive rules having one or more premises and axiom rules with no premises. Each piece of syntax has at least one corresponding rule for \rightarrow , giving the semantics of each syntactic construct.

³Viewed set theoretically, \rightarrow is a ternary relation, i.e. $(\rightarrow) \subset P \times l \times P$.

Actions The first simple reduction rule is for action prefixes:

$$\text{action} \frac{}{\alpha.P \xrightarrow{\alpha} P}$$

We see that a process with an action a prefixing another process P can reduce by emitting that action and then becoming the remaining processes P . e.g.

$$\text{action} \frac{}{a.b.\mathbf{0} \xrightarrow{a} b.\mathbf{0}} \quad (\text{example})$$

Definition 1 (Trace). Labelled transitions can be composed together to form a *trace*, which comprises a sequence of reductions steps, e.g.,

$$a.b.\mathbf{0} \xrightarrow{a} b.\mathbf{0} \xrightarrow{b} \mathbf{0}$$

Each step here is a reduction that needs a *derivation*. We will see below that we can derive reductions for more complex processes by plugging together inference rules.

Parallel composition A parallel composition of processes $P \mid Q$ has two possible reductions given by the following inductive rules:

$$\text{par1} \frac{P \xrightarrow{l} P'}{P \mid Q \xrightarrow{l} P' \mid Q} \quad \text{par2} \frac{Q \xrightarrow{l} Q'}{P \mid Q \xrightarrow{l} P \mid Q'}$$

The *par1* rule says that if we have a process P that can reduce to P' with label l then we can place P in parallel with Q (i.e., $P \mid Q$) and reduce to $P' \mid Q$ with label l . Essentially we are allowing a process within a parallel composition to make some progress. The *par2* provides the symmetric case.

Example 1. We can compose the *action* and *par1* rules to get the following derivation of a single reduction step for the process $a.P \mid b.Q$:

$$\text{par1} \frac{\text{action} \frac{}{a.P \xrightarrow{a} P}}{a.P \mid b.Q \xrightarrow{a} P \mid b.Q} \quad (\text{example})$$

Thus, we have reduced on the left-hand side of the parallel composition, emitting an a action to get the resulting process $P \mid b.Q$. This process could then reduce by applying *par2*:

$$\text{par2} \frac{\text{action} \frac{}{b.Q \xrightarrow{b} Q}}{P \mid b.Q \xrightarrow{b} P \mid Q} \quad (\text{example})$$

These two reduction steps, once derived, can be put together to form the following trace:

$$a.P \mid b.Q \xrightarrow{a} P \mid b.Q \xrightarrow{b} P \mid Q$$

This is not the only possible reduction sequence for the process $(a.P \mid b.Q)$. Another possibility is that we could reduce on the right first, then on the left, giving the trace:

$$a.P \mid b.Q \xrightarrow{b} a.P \mid Q \xrightarrow{a} P \mid Q$$

We can put these two traces together into a *transition graph* that shows the possible traces:

$$\begin{array}{ccc} a.P \mid b.Q & \xrightarrow{a} & P \mid b.Q \\ \downarrow b & & \downarrow a \\ a.P \mid Q & \xrightarrow{a} & P \mid Q \end{array}$$

Note that this represents a *confluent* reduction, i.e., we take different paths in the reduction but end in the same place.

A further reduction is possible for a parallel composition of processes, but in a restricted setting where two parallel processes reduce by dual actions:

$$\text{handshake} \frac{P \xrightarrow{a} P' \quad Q \xrightarrow{\bar{a}} Q'}{P \mid Q \xrightarrow{\tau} P' \mid Q'}$$

In this case we have the processes P and Q making a reduction step *at the same time* to become $P' \mid Q'$ if they reduce by dual actions a and \bar{a} . We can use this to model notions of synchronisation, co-ordination, or communication between processes.

Example 2. Consider the following processes representing a vending machine V and a customer C :

$$\begin{aligned} V &\stackrel{def}{=} \text{coin}.\overline{\text{tea}}.V \\ C &\stackrel{def}{=} \overline{\text{coin}}.\text{tea}.C \end{aligned}$$

Since V can reduce by emitting the action `coin` and C can reduce by emitting the action $\overline{\text{coin}}$ we can get the reduction:

$$\text{handshake} \frac{V \xrightarrow{\text{coin}} \overline{\text{tea}}.V \quad C \xrightarrow{\overline{\text{coin}}} \text{tea}.C}{V \mid C \xrightarrow{\tau} (\overline{\text{tea}}.V \mid \text{tea}.C)} \quad (\text{example})$$

Note that the resulting process can further do a handshake because of the dual `tea` actions, bringing us back to the process $V \mid C$.

Choice A non-deterministic choice between two processes $P + Q$ reduces by reducing just one side, resulting in a process that corresponds to just one side of the choice, given by the rules:

$$\text{choice1} \frac{P_1 \xrightarrow{l} Q}{P_1 + P_2 \xrightarrow{l} Q} \quad \text{choice2} \frac{P_2 \xrightarrow{l} Q}{P_1 + P_2 \xrightarrow{l} Q}$$

Example 3. For a process $b.P + c.Q$, the top-level syntactic construct is non-deterministic choice, so a reduction must necessarily use *choice1* or *choice2*. For example:

$$\text{choice2} \frac{\text{action} \frac{c.Q \xrightarrow{c} Q}{b.P + c.Q \xrightarrow{c} Q}}{b.P + c.Q \xrightarrow{c} Q} \quad (\text{example})$$

Note that when we make a choice, we lose the other branch— above, $b.P$ disappears in the resulting process term. The lectures have more examples.

Restriction The notion of restricting a name for a process, i.e. $\nu a.P$ can be thought of a bit like binding a name a in the scope of P so that the name is “local” to P and cannot be observed outside of it. Its semantics is given by:

$$\text{restriction} \frac{P \xrightarrow{l} P'}{\nu a.P \xrightarrow{l} \nu a.P'} \quad l \neq a \wedge l \neq \bar{a}$$

Here we are saying that we can observe the action given by label l only if it is not the bound name a or its dual \bar{a} . We can use this to enforce handshaking of dual actions.

Example 4. Consider the processes $a.P$ and $\bar{a}.Q$. Their parallel composition has three possible traces:

$$\begin{aligned} (a.P \mid \bar{a}.Q) &\xrightarrow{a} (P \mid \bar{a}.Q) \xrightarrow{\bar{a}} (P \mid Q) \\ (a.P \mid \bar{a}.Q) &\xrightarrow{\bar{a}} (a.P \mid Q) \xrightarrow{a} (P \mid Q) \\ (a.P \mid \bar{a}.Q) &\xrightarrow{\tau} (P \mid Q) \end{aligned}$$

By restricting the process on action a (i.e., the process term $\nu a.(a.P \mid \bar{a}.Q)$) we can make sure that the handshaking trace is the only one possible, preventing the other two traces, with the derivation:

$$\text{restriction} \frac{\text{handshake} \frac{\text{action} \frac{a.P \xrightarrow{a} P}{a.P \xrightarrow{a} P} \quad \text{action} \frac{\bar{a}.Q \xrightarrow{\bar{a}} Q}{\bar{a}.Q \xrightarrow{\bar{a}} Q}}{(a.P \mid \bar{a}.Q) \xrightarrow{\tau} P \mid Q}}{\nu a.(a.P \mid \bar{a}.Q) \xrightarrow{\tau} \nu a.(P \mid Q)} \quad (\text{example})$$

The side condition for *restriction* is satisfied since $a \neq \tau$ and $\bar{a} \neq \tau$. No other derivation is possible. For example, the following is not a valid derivation because it would violate the side condition of *restriction* (that the label does not match the name being restricted over):

$$\text{X} \quad \text{restriction} \frac{\text{par} \frac{\text{action} \frac{a.P \xrightarrow{a} P}{a.P \xrightarrow{a} P}}{(a.P \mid \bar{a}.Q) \xrightarrow{a} (P \mid \bar{a}.Q)}}{\nu a.(a.P \mid \bar{a}.Q) \xrightarrow{a} \nu a.(P \mid \bar{a}.Q)} \quad (\text{example})$$

We can encapsulate the above notion via a theorem:

Theorem 1. Restriction of a name forces handshaking on dual names. That is:

$$\begin{aligned} \forall P, P', Q, Q', a. \quad & P \xrightarrow{a} P' \wedge Q \xrightarrow{\bar{a}} Q' \wedge P \neq P' \wedge Q \neq Q' \\ \Leftrightarrow & \nu a.P \mid Q \xrightarrow{\tau} \nu a.P' \mid Q' \end{aligned}$$

That is, if two process can reduce by emitting dual names a and \bar{a} , then their parallel composition under a restriction of a can only reduce by a handshake, and vice versa.⁴

Finally, a common syntactic shorthand for a several name restrictions is to give a comma-separated list of the restricted names, e.g. $\nu a.\nu b.P$ is written as $\nu a, b.P$.

Process identifiers A process identifier A can reduce given a definition for A which itself can reduce. This is given by the rule:

$$\text{def} \frac{P \xrightarrow{l} Q}{A \xrightarrow{l} Q} \quad (A \stackrel{\text{def}}{=} P)$$

Example 5. For example, if we have $A \stackrel{\text{def}}{=} a.b.A$ then we get the derivation:

$$\text{def} \frac{\text{action} \frac{a.b.A \xrightarrow{a} b.A}{a.b.A \xrightarrow{a} b.A}}{A \xrightarrow{a} b.A} \quad (\text{example})$$

⁴The lecture stated this slightly differently with:

$$\begin{aligned} \forall P, P', Q, Q', a. \quad & P \mid Q \xrightarrow{a} P' \mid Q \xrightarrow{\bar{a}} P' \mid Q' \\ \vee & P \mid Q \xrightarrow{\bar{a}} P \mid Q' \xrightarrow{a} P' \mid Q' \\ \Leftrightarrow & \nu a.P \mid Q \xrightarrow{\tau} \nu a.P' \mid Q' \end{aligned}$$

This is essentially the same as the theorem here, but more complex than necessary. Note that we also make it clear here that P and P' must be distinct in order to get the bi-implication.

Relabelling Relabelling give us a way to rename an emitted action. That is $P[b/a]$ means that for process P , any emitted action a is renamed to the action b . The semantics is given by the reduction:

$$\text{relabel} \frac{P \xrightarrow{l} Q}{P[b/a] \xrightarrow{l[b/a]} P[b/a]}$$

where relabelling $l[b/a]$ means to relabel name a to b for the label l , defined:

$$l[b/a] = \begin{cases} b & \text{when } l = a \\ \bar{b} & \text{when } l = \bar{a} \\ l & \text{otherwise} \end{cases}$$

Relabelling is useful when reusing definitions for different purposes (abstraction and reuse), but we do not make frequent use of it in the course.

2.1 Collected rules

$$\begin{array}{c} \text{action} \frac{}{\alpha.P \xrightarrow{\alpha} P} \\[10pt] \text{par1} \frac{P \xrightarrow{l} P'}{P \mid Q \xrightarrow{l} P' \mid Q} \quad \text{par2} \frac{Q \xrightarrow{l} Q'}{P \mid Q \xrightarrow{l} P \mid Q'} \\[10pt] \text{handshake} \frac{P \xrightarrow{a} P' \quad Q \xrightarrow{\bar{a}} Q'}{P \mid Q \xrightarrow{\tau} P' \mid Q'} \quad \text{restriction} \frac{P \xrightarrow{l} P'}{\nu a.P \xrightarrow{l} \nu a.P'} \quad l \neq a \wedge l \neq \bar{a} \\[10pt] \text{choice1} \frac{P_1 \xrightarrow{l} Q}{P_1 + P_2 \xrightarrow{l} Q} \quad \text{choice2} \frac{P_2 \xrightarrow{l} Q}{P_1 + P_2 \xrightarrow{l} Q} \\[10pt] \text{def} \frac{P \xrightarrow{l} Q}{A \xrightarrow{l} Q} \quad (A \stackrel{\text{def}}{=} P) \quad \text{relabel} \frac{P \xrightarrow{l} Q}{P[b/a] \xrightarrow{l[b/a]} P[b/a]} \end{array}$$

where relabelling $l[b/a]$ means to relabel name a to b for the label l , defined:

$$l[b/a] = \begin{cases} b & \text{when } l = a \\ \bar{b} & \text{when } l = \bar{a} \\ l & \text{otherwise} \end{cases}$$

Example 6. Here is an extended example of a derivation of one possible reduction for the process defined $A \stackrel{\text{def}}{=} a.0 \mid (b.P + c.Q)$:

$$\text{def} \frac{\text{par2} \frac{\text{choice2} \frac{\text{action} \frac{}{c.Q \xrightarrow{c} Q}}{b.P + c.Q \xrightarrow{c} Q}}{a.0 \mid (b.P + c.Q) \xrightarrow{c} a.0 \mid Q}}{A \xrightarrow{c} a.0 \mid Q} \quad (\text{example})$$

3 Modelling locks, semaphores, and race conditions

We can use CCS to model common concurrent programming constructs, helping us to understand their behaviour. The following accompanies the lectures for this concept.

The following is a simple CCS model for the vending machine of `TeaRoomInitial.java`:

$$\begin{aligned}
 (\text{accounting}) \quad & \text{Inc}(p).P \stackrel{\text{def}}{=} r(p).\overline{w}(p+1).P \\
 & \text{Dec}(s).P \stackrel{\text{def}}{=} r(s).\overline{w}(s-1).P \\
 & A.P \stackrel{\text{def}}{=} \text{Dec}(s).\text{Inc}(p).P \\
 (\text{vend}) \quad & V \stackrel{\text{def}}{=} \text{coin}.\overline{\text{lock}}.A.\overline{\text{unlock}}.\overline{\text{tea}} \\
 (\text{machine}) \quad & M \stackrel{\text{def}}{=} V \mid M
 \end{aligned}$$

The “accounting” agent A provides the behaviour of reading and then updating two variables: incrementing the profit p and decrementing the supply s . Note that we abuse notation slightly here and definition *parametric* processes $\text{Inc}(p)$ and $\text{Dec}(s)$ parameterised by the name of some variable we are changing. This is not real CCS syntax, but a bit of meta-level syntactic sugar to simplify the model.

The vending machine receives a `coin`, then requests the lock before trying to behave as an accounting agent, then unlocking and sending tea. Since a machine is a shared object for which the `vend` method can be called arbitrarily and concurrently, a machine M is an infinite parallel composition of V .

A customer and a lock (a shared object) are defined:

$$\begin{aligned}
 (\text{lock}) \quad & L \stackrel{\text{def}}{=} \text{lock}.\overline{\text{unlock}}.L \\
 (\text{customer}) \quad & C \stackrel{\text{def}}{=} \overline{\text{coin}}.\text{tea}.C
 \end{aligned}$$

The following CCS agent then models two consumers at the vending machine:

$$\nu\text{coin}.\nu\text{tea}.\nu\text{lock}.\nu\text{unlock}.(C \mid C \mid M \mid L)$$

We are *restricting* the names `coin`, `tea`, `lock` and `unlock` so that handshaking is guaranteed (see Theorem 1). The following gives a trace of the model, but we elide the restriction since it is not necessary to get the following trace.

A process P is highlighted as \overline{P} when we are going to expand its definition in the next line, where P marks the expanded definition. Pairs of dual (complementary) actions are highlighted as \overline{a} and a when they are about to handshake via a τ step.

$$\begin{aligned}
 \overline{C} \mid C \mid \overline{M} \mid L & \equiv \overline{\text{coin}}.\text{tea}.C \mid C \mid \overline{\text{coin}}.\overline{\text{lock}}.A.\overline{\text{unlock}}.\overline{\text{tea}} \mid M \mid L \\
 & \xrightarrow{\tau} \text{tea}.C \mid C \mid \overline{\text{lock}}.A.\overline{\text{unlock}}.\overline{\text{tea}} \mid M \mid \overline{L} \\
 & \equiv \text{tea}.C \mid C \mid \overline{\text{lock}}.A.\overline{\text{unlock}}.\overline{\text{tea}} \mid M \mid \overline{\text{lock}}.\overline{\text{unlock}}.L \\
 & \xrightarrow{\tau} \text{tea}.C \mid \overline{C} \mid A.\overline{\text{unlock}}.\overline{\text{tea}} \mid \overline{M} \mid \overline{\text{unlock}}.L \\
 & \equiv \text{tea}.C \mid \overline{\text{coin}}.\text{tea}.C \mid A.\overline{\text{unlock}}.\overline{\text{tea}} \mid \overline{\text{coin}}.\overline{\text{lock}}.A.\overline{\text{unlock}}.\overline{\text{tea}} \mid M \mid \overline{\text{unlock}}.L \\
 & \xrightarrow{\tau} \text{tea}.C \mid \text{tea}.C \mid \underbrace{A.\overline{\text{unlock}}.\overline{\text{tea}}}_{\text{vend 1}} \mid \underbrace{\overline{\text{lock}}.A.\overline{\text{unlock}}.\overline{\text{tea}}}_{\text{vend 2}} \mid M \mid \overline{\text{unlock}}.L
 \end{aligned}$$

The second vending process is not able to acquire the lock since there is no dual lock action that is the head prefix of any other process. The “vend 2” process has to wait for the first one (“vend 1”)

to unlock via $\overline{\text{unlock}}$ which then puts the lock L back to its initial configuration. That is, after some steps dealing with accounting (which is now atomic) we get:

$$\begin{aligned} & \text{tea}.C \mid \text{tea}.C \mid \underbrace{\overline{\text{tea}}}_{\text{vend 1}} \mid \underbrace{\overline{\text{lock}.A.\overline{\text{unlock}}.\text{tea}}}_{\text{vend 2}} \mid M \mid \overline{L} \\ \equiv & \text{tea}.C \mid \text{tea}.C \mid \overline{\text{tea}} \mid \overline{\text{lock}}.A.\overline{\text{unlock}}.\overline{\text{tea}} \mid M \mid \overline{\text{lock}}.\overline{\text{unlock}}.L \end{aligned}$$

We define the following CCS model of `TeaRoomSem.java` which reuses the accounting agent A :


$$\begin{aligned} (1\text{-semaphore}) \quad S_1 &\stackrel{\text{def}}{=} \text{acq}.\text{rel}.S_1 \\ (n\text{-semaphore}) \quad S_n &\stackrel{\text{def}}{=} S_1 \mid S_{n-1} \\ (\text{vend}) \quad V &\stackrel{\text{def}}{=} \text{coin}.\overline{\text{acq}}.A.\overline{\text{rel}}.\overline{\text{tea}} \\ (\text{machine}) \quad M &\stackrel{\text{def}}{=} V \mid M \\ (\text{customer}) \quad C &\stackrel{\text{def}}{=} \overline{\text{coin}}.\text{tea}.C \end{aligned}$$

A semaphore with n permissions (capacity) is modelled by S_n which composes in parallel n copies of the binary semaphore agent S_1 . We can see that a binary semaphore is basically the same as our previous model of a lock L . However, now n processes can be in their critical section.

The following trace models two consumers at a vending machine with a semaphore that has two permissions:

$$\begin{aligned} \overline{C} \mid C \mid M \mid S_2 &\equiv \overline{\text{coin}}.\text{tea}.C \mid C \mid \text{coin}.\overline{\text{acq}}.A.\overline{\text{rel}}.\overline{\text{tea}} \mid M \mid S_2 \\ &\xrightarrow{\tau} \text{tea}.C \mid C \mid \overline{\text{acq}}.A.\overline{\text{rel}}.\overline{\text{tea}} \mid M \mid S_2 \\ &\equiv \text{tea}.C \mid C \mid \overline{\text{acq}}.A.\overline{\text{rel}}.\overline{\text{tea}} \mid M \mid \text{acq}.\text{rel}.S_1 \mid S_1 \\ &\xrightarrow{\tau} \text{tea}.C \mid \overline{C} \mid A.\overline{\text{rel}}.\overline{\text{tea}} \mid M \mid \text{rel}.S_1 \mid S_1 \\ &\equiv \text{tea}.C \mid \overline{\text{coin}}.\text{tea}.C \mid A.\overline{\text{rel}}.\overline{\text{tea}} \mid \text{coin}.\overline{\text{acq}}.A.\overline{\text{rel}}.\overline{\text{tea}} \mid M \mid \text{rel}.S_1 \mid S_1 \\ &\xrightarrow{\tau} \text{tea}.C \mid \text{tea}.C \mid A.\overline{\text{rel}}.\overline{\text{tea}} \mid \overline{\text{acq}}.A.\overline{\text{rel}}.\overline{\text{tea}} \mid M \mid \text{rel}.S_1 \mid S_1 \\ &\equiv \text{tea}.C \mid \text{tea}.C \mid A.\overline{\text{rel}}.\overline{\text{tea}} \mid \overline{\text{acq}}.A.\overline{\text{rel}}.\overline{\text{tea}} \mid M \mid \text{rel}.S_1 \mid \text{acq}.\text{rel}.S_1 \\ &\xrightarrow{\tau} \text{tea}.C \mid \text{tea}.C \mid A.\overline{\text{rel}}.\overline{\text{tea}} \mid A.\overline{\text{rel}}.\overline{\text{tea}} \mid M \mid \text{rel}.S_1 \mid \text{rel}.S_1 \end{aligned}$$

We now have a race condition in the reduction of the yellow highlighted sub-process, which we show separate from the rest of the context (which is independent) setting $Q = \overline{\text{rel}}.\overline{\text{tea}}$:



Note that the middle term is achieved when the left and right hand processes interleave their reads, leading to a state where both processes have read the same value of s (the supply) and will write the same value of $s - 1$, shadowing the fact that **two** processes are actually reducing the supply.

Furthermore, this can hide the situation where the supply is only one ($s = 1$) and therefore only one process is allowed to consume from the supply. But if both processes interleave their reads as in the above, they will both see a view of the world where $s = 1$ and erroneously assume there is enough in the supply for them both to proceed.

Thus we see that semaphores alone cannot prevent *race conditions*. If semaphore-guarded code contains shared state, we must also enforce mutual exclusion on that shared state via an additional binary semaphore or mutual-exclusion lock.

4 Structural congruence

A basic notion of equality can be ascribed to processes which explains when processes are treated as being the same. This equality relation, which we write as \equiv , is known as *structural congruence*. Which we define in chunks here based on the main syntactic construct of interest.

Firstly, for parallel composition, structural congruence is defined:

$$\begin{aligned} P \mid Q &\equiv Q \mid P \\ (P \mid Q) \mid R &\equiv P \mid (Q \mid R) \\ P \mid \mathbf{0} &\equiv P \end{aligned}$$

i.e., parallel composition is commutative and associative, and has the inactive process $\mathbf{0}$ as a unit; parallel composition is a commutative monoid with $\mathbf{0}$.

For choice, we get the equations:

$$\begin{aligned} P + Q &\equiv Q + P \\ (P + Q) + R &\equiv P + (Q + R) \\ P + \mathbf{0} &\equiv P \\ P + P &\equiv P \end{aligned}$$

Thus, choice is commutative, associative, has $\mathbf{0}$ as a unit, and is also idempotent: choosing between the same thing is no choice at all.

Finally, for restriction we have some more involved rules with side conditions, that related to how we can distribute bindings (that is, move them with respect to other pieces of syntax). In these equations we make use of a meta-level operation on processes $\text{FN}(P)$ which returns the set of free names (i.e. unrestricted names) in P . For example $\text{FN}(a.b.\mathbf{0}) = \{a, b\}$ but $\text{FN}(a.\mathbf{0} \mid \nu b.(b.\mathbf{0})) = \{a\}$.

$$\begin{aligned} \nu a.P &\equiv P && \text{if } a, \bar{a} \notin \text{FN}(P) \\ \nu a.(P + Q) &\equiv \nu a.P + \nu a.Q \\ \nu a.(P \mid Q) &\equiv \nu a.P \mid \nu a.Q && \text{if } (a \in \text{FN}(P) \Rightarrow \bar{a} \notin \text{FN}(Q)) \wedge (a \in \text{FN}(Q) \Rightarrow \bar{a} \notin \text{FN}(P)) \\ \nu a.(P \mid Q) &\equiv (\nu a.P) \mid Q && \text{if } a, \bar{a} \notin \text{FN}(Q) \end{aligned}$$

The first allows us to drop redundant restrictions. The second allows us to distribute restriction inside choice. The third allows us to distribute restriction inside parallel composition only if P and Q cannot handshake via a . The fourth equation states that we can shrink or grow the scope of a restriction over processes that don't mention the restricted name.

We can apply structural congruences to simplify processes. For example:

$$P + (\mathbf{0} \mid P) \equiv P \quad \text{(example)}$$

5 Bisimulation

Often we want to be able to reason about whether two processes behave the same. For example, when programming, we might come up with a refactoring or simplification which we want to ensure behaves in the same way as the old program. Or, we might one to use one process as a specification for another, giving a simple but inefficient definition as the specification, against which we want to compare the complicated but efficient implementation that we use. We can describe process equivalence by comparing process behaviour: that is, do the two processes behave in the same way?

The notion of behaviour equivalence we study in this course is called *strong bisimulation*. This notion of behaviour equivalence compares step-by-step what two processes can do. Roughly, if we say that a process P is *bisimilar* to a process Q one process if P can make a reduction step with label l then so can Q , and vice versa, and the processes reached after reduction are also bisimilar. This is encapsulated by the following definition, which is actually a property of a relation.

Definition 2 (Bisimulation). A relation \mathcal{R} between processes is a bisimulation relation iff it satisfies the following: for all processes P and Q then:

$$\begin{aligned} P \mathcal{R} Q \quad &\Rightarrow \forall P', l. \ P \xrightarrow{l} P' \Rightarrow (\exists Q''. Q \xrightarrow{l} Q'' \wedge P' \mathcal{R} Q'') \\ &\wedge \forall Q', l. \ Q \xrightarrow{l} Q' \Rightarrow (\exists P''. P \xrightarrow{l} P'' \wedge P'' \mathcal{R} Q') \end{aligned}$$

Definition 3 (Strong bisimulation). Two processes P and Q are *bisimilar* if $P \sim Q$ where \sim is the largest bisimulation, defined:

$$\sim = \bigcup \{ \mathcal{R} \mid \mathcal{R} \text{ is a bisimulation} \}$$

(i.e. \sim incorporates all possible bisimulation relations over processes).

The following are some small examples:

- $(a.0 \mid \bar{a}.0) \sim a.\bar{a}.0 + \bar{a}.a.0 + \tau.0$;
- $P \sim a.Q$ where $P \stackrel{def}{=} a.b.P$ and $Q \stackrel{def}{=} b.a.Q$;
- $a.(b + c) \not\sim a.b + a.c$ (i.e., distributing a prefix inside a choice leads to different behaviour).

The lecture gives more examples and motivates this definition further.